12/3/2019 PrairieLearn

HW5.1. Spam Detection

Bayesian Spam Filtering

A common way to detect spam in emails is to use a *Naive Bayes classifier*, which identifies key words that are frequently used in spam messages and then calculates the probability that a given message is either spam or not spam using the words contained in the message.

If you are interested, you can read more about <u>Bayes' theorem</u> and <u>Naive</u> <u>Bayes spam filtering</u>, but you do not need to understand it to do this problem.

As always, feel free to work on this incrementally to check your progress. You have unlimited attempts for this problem.

1. You are given a dataset of sample messages as spam. csv, using the Pandas library read this file into a DataFrame named df. This dataset gives text messages and their classifications, either spam or ham (not spam).

Here is part of the dataset so you can get an idea of the values:

	Label	Text
0	ham	Go until jurong point, crazy Available only in bugis n great world la e buffet Cine there got amore wat
1	ham	Ok lar Joking wif u oni
2	spam	Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to receive entry question(std txt rate)T&C's apply 08452810075over18's
3	ham	U dun say so early hor U c already then say
4	ham	Nah I don't think he goes to usf, he lives around here though

5 rows x 2 columns

2. At this point, lets split the dataset into spam and ham values. Save the rows that have only spam or ham into the variables spam and ham.

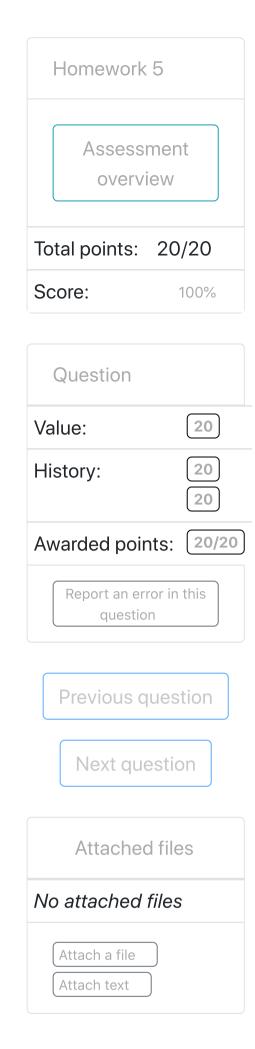
We can use this information to approximate the overall probability of receiving a spam or ham message, P(S) and P(H), using:

 $\frac{\text{number of messages}}{\text{total messages}}$

Calculate these values and store as prob_spam and prob_ham.

3. We will want to decompose a message into a set of words, or *tokens*. Implement the function tokenize_text below, which should take as input a string and return a set of tokens. Use NLTK's <u>word_tokenize()</u> to split a string into a list of words. Then, remove anything that is not a word (str.isalpha() may be helpful).

To help identify words, we can reduce them to their *lemma*, which is the base form of the word. For example, the words "produced" and "production" have the same lemma, "produce", this will reduce the



12/3/2019 PrairieLearn

number of redundant entries we get. For this we can use NLTK's WordNetLemmatizer; replace all tokens with their lower-cased lemma.

```
# Example code
lemmatizer = WordNetLemmatizer()
lemma = lemmatizer.lemmatize("rock")
```

4. Now, we can calculate the probability of finding a word in a spam or a ham message. This is equivalent to finding the percentage that each word occurs in either. Loop through each word in each message, and if it is spam increase the word's spam probability by 1/N (where N is the total number of messages), and if it is ham increase the word's ham probability by the same amount. This is a similar method to generating the probabilities in Markov chains in the lesson activity.

Create the two dictionaries, prob_spam_words and prob_ham_words to store the probability of each word. Each dictionary should map a string token to its probability of showing up in a spam or ham message. You can do this by looping over the tokens in each message and either incrementing the probability by 1/N if the word exists in the dictionary, or setting the probability to 1/N if it does not exist yet.

5. After generating all the above probabilities, we now have all the information needed to classify a piece of text as spam or not.

You are given the function <code>get_class_probability()</code> that will give you the probability of getting a body of text given it is ham or spam, $P\left(\mathbf{W}\mid S\right)$ and $P\left(\mathbf{W}\mid H\right)$. Using Bayes' theorem, we can turn this into the probability that some piece of text is spam or not:

$$P(S \mid \mathbf{W}) \propto P(\mathbf{W} \mid S) P(S)$$

 $P(H \mid \mathbf{W}) \propto P(\mathbf{W} \mid H) P(H)$

(Because we are comparing both probabilities against each other, we can drop the denominators since they're the same.)

To find the probability a message is spam, we can multiply the probability given by get_class_probability() by the overall probability of spam. You can get a similar result for the probability of ham. To classify a message, we can just check which probability of the two is higher. I.e. if the spam probability is larger, then the message is classified as spam.

You are given an example named test_message to classify. Tokenize the message and find the probability that it is spam and ham, save these as prob_msg_spam and prob_msg_ham. Finally, determine if it is spam or not and save this result in is_spam.

You are given the following variable:

Name	Туре	Description
test_message	string	Text to classify

12/3/2019 PrairieLearn

Your code snippet should define the following variables:

Name	Туре	Description	
df	pandas dataframe	Full dataset	
tokenize_te xt	function	Function to clean a message and split it into tokens	
spam	pandas dataframe	Dataset of only spam messages	
ham	pandas dataframe	Dataset of only ham messages	
prob_spam	decimal number	Probability of having a spam message	
prob_ham	decimal number	Probability of having a ham message	
prob_spam_w ords	dictionary	Probability of having a specific word, given a spam message	
prob_ham_wo rds	dictionary	Probability of having a specific word, given a ham message	
prob_msg_sp am	decimal number	Probability that the test message is spam	
prob_msg_ha m	decimal number	Probability that the test message is ham	
is_spam	boolean	Is the test message spam?	

user_code.py

```
1 import numpy as np
 <sup>2</sup> import pandas as pd
 3 import nltk
    from nltk.corpus import stopwords
    from nltk.stem import WordNetLemmatizer
    # 1. Load dataset
    df = pd.read_csv("spam.csv")
10
    # 2. Split Spam and Ham
12
    spam = df[df['Label'] == 'spam']
14
    ham = df[df['Label'] == 'ham']
15
16
    num_total = len(spam.to_numpy()) + len(ham.to_numpy())
17
    num_spam = len(spam.to_numpy())
18
    num_ham = len(ham.to_numpy())
19
20
    prob_spam = float(num_spam)/num_total
    prob_ham = float(num_ham)/num_total
22
23
24
25
26
    # 3. Create tokenizer function
27
```

```
28
    def tokenize_text(message):
29
        0.00
30
        message: A string containing a single sentence.
31
        returns: A list of tokens.
32
33
        # Implement me!
34
        # - Tokenize the message
35
        # - Remove non word tokens and stopwords
36
        # - Convert tokens to lemmas
37
        # - Convert everything to lowercase
38
        stopWords = set(stopwords.words('english'))
39
        message = nltk.word_tokenize(message)
40
        ret = []
41
42
        for word in message:
43
44
            if str.isalpha(word) and word not in stopWords:
45
46
                lemmatizer = WordNetLemmatizer()
47
                lemma = lemmatizer.lemmatize(word)
48
                lemma = lemma.lower()
49
                ret.append(lemma)
50
51
        return ret
52
# an example to test your function:
54 print(tokenize_text(df.values[0, 1]))
55 # the expected output:
56 print(['go', 'jurong', 'point', 'available', 'bugis', 'n',
        'great', 'world', 'la', 'e', 'buffet', 'cine', 'got',
        'amore', 'wat'])
57
58
    # 4. Get Word Probabilities
59
60
    prob\_spam\_words = \{\}
61
    prob_ham_words = {}
62
63
64
    # Uncomment this once you get to it.
65
66
    for text in spam.values[:,1]:
67
        tokens = tokenize_text(text)
68
        for token in tokens:
69
            if token not in prob_spam_words:
70
                prob_spam_words[token] = 1./num_total
71
            else:
72
                prob_spam_words[token] += 1./num_total
73
    for text in ham.values[:,1]:
75
        tokens = tokenize_text(text)
76
        for token in tokens:
77
            if token not in prob_ham_words:
78
                prob_ham_words[token] = 1./num_total
79
            else:
80
                prob_ham_words[token] += 1./num_total
81
82
83
    # 5. Classify message
84
85
    def get_class_probability(tokens, probabilities):
86
        eta = 0
87
        for token in tokens:
88
            p = 1e-4
```

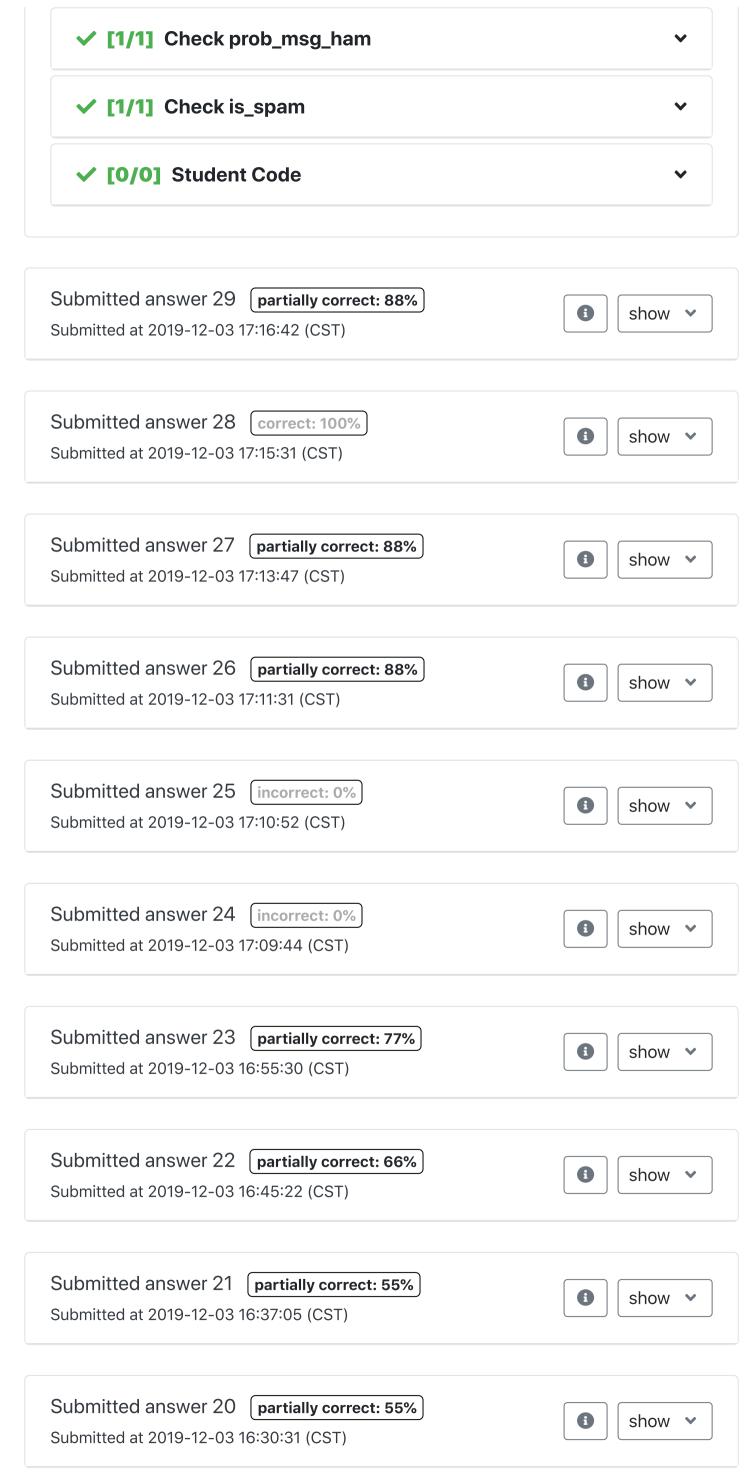
12/3/2019 PrairieLearn

```
89
               if token in probabilities:
   90
                   p = probabilities[token]
   91
               eta += np.log(p)
   92
           return np.exp(eta)
   93
   94
   95
   96
       pws = get_class_probability(test_message,prob_spam_words)
   97
       prob_msg_spam = pws*prob_spam
   98
       phs = get_class_probability(test_message,prob_ham_words)
       prob_msg_ham = phs*prob_ham
  100
       is_spam = False if prob_msg_spam > prob_msg_ham else True
  101
                                                               Restore original file
                     Save only
Save & Grade
```

Submitted answer 30 correct: 100% hide ^ Submitted at 2019-12-03 17:17:52 (CST) Score: 9/9 (100%) Output ['go', 'jurong', 'point', 'available', 'bugis', 'n', 'great', 'world', 'la', 'e', 'buffet', 'cine', 'got', 'amore', 'wat'] ['go', 'jurong', 'point', 'available', 'bugis', 'n', 'great', 'world', 'la', 'e', 'buffet', 'cine', 'got', 'amore', 'wat']

Test Results

✓ [1/1] Check df ✓ [1/1] Check tokenize_text **√** [1/1] Check spam **√** [1/1] Check ham ✓ [1/1] Check prob_spam ✓ [1/1] Check prob_ham ✓ [1/1] Check prob_spam_words



PrairieLearn Submitted answer 19 partially correct: 55% show Y Submitted at 2019-12-03 16:29:38 (CST) Submitted answer 18 partially correct: 55% show 🕶 Submitted at 2019-12-03 16:25:16 (CST) Submitted answer 17 partially correct: 55% show 🕶 Submitted at 2019-12-03 16:24:13 (CST) Submitted answer 16 | partially correct: 55% show 🕶 Submitted at 2019-12-03 16:15:46 (CST) Submitted answer 15 partially correct: 55% show ~ Submitted at 2019-12-03 16:14:15 (CST) Submitted answer 14 saved, not graded show ~ Submitted at 2019-12-03 16:13:19 (CST) Submitted answer 13 incorrect: 0% show ~ Submitted at 2019-12-03 12:15:35 (CST) Submitted answer 12 partially correct: 55% show ~ Submitted at 2019-12-03 12:12:12 (CST) Submitted answer 11 partially correct: 55% show ~ Submitted at 2019-12-03 12:10:22 (CST) Submitted answer 10 partially correct: 55% show ~ Submitted at 2019-12-03 12:09:08 (CST) Submitted answer 9 partially correct: 55% show ~ Submitted at 2019-12-03 12:05:32 (CST) Submitted answer 8 | partially correct: 33% show 💙 Submitted at 2019-12-03 12:04:49 (CST)

Submitted at 2019-12-03 12:03:27 (CST)

PrairieLearn

SI IOW

Submitted answer 6 partially correct: 11%

Submitted at 2019-12-03 12:02:20 (CST)

Submitted answer 5 partially correct: 11%

Submitted at 2019-12-03 11:58:19 (CST)

show >

show 🕶

Submitted answer 4 partially correct: 11%

Submitted at 2019-12-03 11:56:58 (CST)

show •

Submitted answer 3 partially correct: 11%

Submitted at 2019-12-03 11:54:00 (CST)

show v

Submitted answer 2 [incorrect: 0%]

Submitted at 2019-12-03 11:53:25 (CST)

show >

Submitted answer 1 partially correct: 11%

Submitted at 2019-12-03 11:48:56 (CST)

show >